



Feature Vector $\phi(x) = [\phi_1(x), \dots, \phi_d(x)]$.

$$\mathbf{w} \cdot \phi(x) = \sum_{j=1}^d w_j \phi_j(x)$$

Linear Classifier

$$f_w(x) = \text{sign}(\mathbf{w} \cdot \phi(x)) = \begin{cases} +1 & \text{if } \mathbf{w} \cdot \phi(x) > 0 \\ -1 & \text{if } \mathbf{w} \cdot \phi(x) < 0 \\ ? & \text{if } \mathbf{w} \cdot \phi(x) = 0 \end{cases}$$

The score on an example (x, y) is $\mathbf{w} \cdot \phi(x)$, how confident we are in predicting +1.

The margin on an example (x, y) is $(\mathbf{w} \cdot \phi(x))y$, how correct we are.

Zero One Loss for binary classifier

$$\text{Loss}_{0-1}(x, y, \mathbf{w}) = \mathbf{1}[f_w(x) \neq y] \\ = \mathbf{1}[\underbrace{(\mathbf{w} \cdot \phi(x))y}_{\text{margin}} \leq 0]$$

Linear Regression

The residual is $(\mathbf{w} \cdot \phi(x)) - y$, the amount by which prediction $f_w(x) = \mathbf{w} \cdot \phi(x)$ overshoots the target y .

Squared Loss and absdev

$$\text{Loss}_{\text{squared}}(x, y, \mathbf{w}) = \underbrace{(f_w(x) - y)^2}_{\text{residual}}$$

$$\text{Loss}_{\text{absdev}}(x, y, \mathbf{w}) = |\mathbf{w} \cdot \phi(x) - y|$$

Objective : Minimize the Train Loss

$$\text{TrainLoss}(\mathbf{w}) = \frac{1}{|\mathcal{D}_{\text{train}}|} \sum_{(x,y) \in \mathcal{D}_{\text{train}}} \text{Loss}(x, y, \mathbf{w}) \\ \min_{\mathbf{w} \in \mathbb{R}^d} \text{TrainLoss}(\mathbf{w})$$

Gradient Descent

Initialize $\mathbf{w} = [0, \dots, 0]$

For $t = 1, \dots, T$:

$$\mathbf{w} \leftarrow \mathbf{w} - \underbrace{\eta}_{\text{step size}} \underbrace{\nabla_{\mathbf{w}} \text{TrainLoss}(\mathbf{w})}_{\text{gradient}} \\ \text{TrainLoss}(\mathbf{w}) = \frac{1}{|\mathcal{D}_{\text{train}}|} \sum_{(x,y) \in \mathcal{D}_{\text{train}}} (\mathbf{w} \cdot \phi(x) - y)^2$$

$$\nabla_{\mathbf{w}} \text{TrainLoss}(\mathbf{w}) = \frac{1}{|\mathcal{D}_{\text{train}}|} \sum_{(x,y) \in \mathcal{D}_{\text{train}}} 2(\underbrace{\mathbf{w} \cdot \phi(x) - y}_{\text{prediction} - \text{target}})\phi(x)$$

Stochastic Gradient Descent

For each $(x, y) \in \mathcal{D}_{\text{train}}$:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{Loss}(x, y, \mathbf{w})$$

Hinge Loss


$$\text{Loss}_{\text{hinge}}(x, y, \mathbf{w}) = \max\{1 - (\mathbf{w} \cdot \phi(x))y, 0\}$$

Gradient of Hinge Loss

$$\nabla_{\mathbf{w}} \text{Loss}_{\text{hinge}}(x, y, \mathbf{w}) = \begin{cases} -\phi(x)y & \text{if } \mathbf{w} \cdot \phi(x)y < 1 \\ 0 & \text{if } \mathbf{w} \cdot \phi(x)y \geq 1 \end{cases}$$

Logistic Regression

$$\text{Loss}_{\text{logistic}}(x, y, \mathbf{w}) = \log(1 + e^{-(\mathbf{w} \cdot \phi(x))y})$$



The logistic function maps $(-\infty, \infty)$ to $[0, 1]$:

$$\sigma(z) = (1 + e^{-z})^{-1}$$

Derivative:

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$



Approximation error: how good is the hypothesis class?

Estimation error: how good is the learned predictor relative to the hypothesis class?

$$\underbrace{\text{Err}(f) - \text{Err}(g)}_{\text{estimation}} + \underbrace{\text{Err}(g) - \text{Err}(f^*)}_{\text{approximation}}$$

Regularization

Initialize $\mathbf{w} = [0, \dots, 0]$

For $t = 1, \dots, T$:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta (\nabla_{\mathbf{w}} [\text{TrainLoss}(\mathbf{w})] + \lambda \mathbf{w})$$

K Means

$$\text{Loss}_{\text{kmeans}}(z, \mu) = \sum_{i=1}^n \|\phi(x_i) - \mu_{z_i}\|^2$$

For each point $i = 1, \dots, n$:

Assign i to cluster with closest centroid:

$$z_i \leftarrow \arg \min_{k=1, \dots, K} \|\phi(x_i) - \mu_k\|^2$$

For each cluster $k = 1, \dots, K$:

Set μ_k to average of points assigned to cluster k :

$$\mu_k \leftarrow \frac{1}{|\{i : z_i = k\}|} \sum_{i: z_i = k} \phi(x_i)$$

Initialize μ_1, \dots, μ_K randomly.

For $t = 1, \dots, T$:

Step 1: set assignments z given μ

Step 2: set centroids μ given z

Legend: b actions/state, solution depth d , maximum depth D

Algorithm	Action costs	Space	Time
DFS	zero	$O(D)$	$O(b^D)$
BFS	constant ≥ 0	$O(b^d)$	$O(b^d)$
DFS-ID	constant ≥ 0	$O(d)$	$O(b^d)$
Backtracking	any	$O(D)$	$O(b^D)$

UCS Algorithm

Add s_{start} to frontier (priority queue)

Repeat until frontier is empty:

Remove s with smallest priority p from frontier

If $\text{IsEnd}(s)$: return solution

Add s to explored

For each action $a \in \text{Actions}(s)$:

Get successor $s' \leftarrow \text{Succ}(s, a)$

If s' already in explored: continue

Update frontier with s' and priority $p + \text{Cost}(s, a)$

DP vs UCS

N total states, n of which are closer than end state

Algorithm	Cycles?	Action costs	Time/space
DP	no	any	$O(N)$
UCS	yes	≥ 0	$O(n \log n)$

Structured Perceptron Algorithm

• For each action: $\mathbf{w}[a] \leftarrow 0$

• For each iteration $t = 1, \dots, T$:

• For each training example $(x, y) \in \mathcal{D}_{\text{train}}$:

• Compute the minimum cost path y' given \mathbf{w}


• For each action $a \in y$: $\mathbf{w}[a] \leftarrow \mathbf{w}[a] - 1$

• For each action $a \in y'$: $\mathbf{w}[a] \leftarrow \mathbf{w}[a] + 1$

A* Algorithm

Run uniform cost search with modified edge costs:

$$\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s)$$



A heuristic h is consistent if

- $\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s) \geq 0$
- $h(s_{\text{end}}) = 0$.

Heuristic is admissible if $h(s) \leq \text{FutureCost}(s)$

Relaxation

A relaxation P_{rel} of a search problem P has costs that satisfy:

$$\text{Cost}_{\text{rel}}(s, a) \leq \text{Cost}(s, a).$$

Given a relaxed search problem P_{rel} , define the relaxed heuristic $h(s) = \text{FutureCost}_{\text{rel}}(s)$, the minimum cost from s to an end state using $\text{Cost}_{\text{rel}}(s, a)$.

And also that means $h(s)$ is consistent.

Note : Heuristic Tradeoff: efficiency(relaxed) vs

tightness(not too relaxed)

Note:if h_1 and h_2 are consistent then $\max(h_1, h_2)$ is too.

MDP



Search State plus Reward and Transition Prob, Discr Factor

A policy π is a mapping from each state $s \in \text{States}$ to an action $a \in \text{Actions}(s)$.


Path: $s_0, a_1, r_1, s_1, a_2, r_2, s_2, \dots$ (action, reward, new state).

The utility with discount γ is

$$u_1 = r_1 + \gamma r_2 + \gamma^2 r_3 + \gamma^3 r_4 + \dots$$

Let $V_{\pi}(s)$ be the expected utility received by following policy π from state s .

Let $Q_{\pi}(s, a)$ be the expected utility of taking action a from state s , and then following policy π .



Policy Evaluation: recurrences and Algorithm

$$V_{\pi}(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ Q_{\pi}(s, \pi(s)) & \text{otherwise.} \end{cases}$$

$$Q_{\pi}(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\pi}(s')]$$

Initialize $V_{\pi}^{(0)}(s) \leftarrow 0$ for all states s .

For iteration $t = 1, \dots, t_{\text{PE}}$:

For each state s :

$$V_{\pi}^{(t)}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [\text{Reward}(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]$$

$$\text{Convergence: } \max_{s \in \text{States}} |V_{\pi}^{(t)}(s) - V_{\pi}^{(t-1)}(s)| \leq \epsilon$$

Optimal value and Policy

The optimal value $V_{\text{opt}}(s)$ is the maximum value attained by any policy.

$$Q_{\text{opt}}(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}(s')]$$

$$V_{\text{opt}}(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} Q_{\text{opt}}(s, a) & \text{otherwise.} \end{cases}$$

$$\pi_{\text{opt}}(s) = \arg \max_{a \in \text{Actions}(s)} Q_{\text{opt}}(s, a)$$

Value Iteration Algorithm:

Initialize $V_{\text{opt}}^{(0)}(s) \leftarrow 0$ for all states s .

For iteration $t = 1, \dots, t_{\text{VI}}$:

For each state s :

$$V_{\text{opt}}^{(t)}(s) \leftarrow \max_{a \in \text{Actions}(s)} \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}^{(t-1)}(s')]$$

Note: converges if mdp is acyclic or disc factor < 1

Reinforcement Learning: no T or R defined in MDP.

Model based Monte Carlo(optimal policy)

Data: $s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$

$$\hat{T}(s, a, s') = \frac{\# \text{ times } (s, a, s') \text{ occurs}}{\# \text{ times } (s, a) \text{ occurs}}$$

$$\widehat{\text{Reward}}(s, a, s') = r \text{ in } (s, a, r, s')$$

$$\hat{Q}_{\text{opt}}(s, a) = \sum_{s'} \hat{T}(s, a, s') [\widehat{\text{Reward}}(s, a, s') + \gamma \hat{V}_{\text{opt}}(s')]$$

Model Free Monte Carlo

$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$

$$u_t = r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \dots$$

$\hat{Q}_{\pi}(s, a)$ = average of u_t where $s_{t-1} = s, a_t = a$
(and s, a doesn't occur in s_0, \dots, s_{t-2})

On each (s, a, u) :

$$\hat{Q}_{\pi}(s, a) \leftarrow \hat{Q}_{\pi}(s, a) - \eta [\underbrace{\hat{Q}_{\pi}(s, a)}_{\text{prediction}} - \underbrace{u}_{\text{target}}]$$

Bootstrapping Methods: SARSA

The main advantage that SARSA offers over model-free Monte Carlo is that we don't have to wait until the end of the episode to update the Q-value.

On each (s, a, r, s', a') :

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta) \hat{Q}_{\pi}(s, a) + \eta [\underbrace{r}_{\text{data}} + \underbrace{\gamma \hat{Q}_{\pi}(s', a')}_{\text{estimate}}]$$

Q-Learning Algorithm(Off Policy and no Succ States reqd)

On each (s, a, r, s') :

$$\hat{Q}_{\text{opt}}(s, a) \leftarrow (1 - \eta) \hat{Q}_{\text{opt}}(s, a) + \eta (\underbrace{r + \gamma \hat{V}_{\text{opt}}(s')}_{\text{prediction}} - \underbrace{\hat{Q}_{\text{opt}}(s, a)}_{\text{target}})$$

Recall: $\hat{V}_{\text{opt}}(s') = \max_{a' \in \text{Actions}(s')} \hat{Q}_{\text{opt}}(s', a')$

Exploration Policy: Exploration vs Exploitation(RL tmplat)

For $T = 1, 2, 3, \dots$

Choose action $a = \pi_{\text{act}}(s_{t-1})$ (how?)

Receive reward r_t and observe new state s_t

Update parameters (how?)

Epsilon Greedy Policy for choosing actions

$$\pi_{\text{act}}(s) = \begin{cases} \arg \max_{a \in \text{Actions}} \hat{Q}_{\text{opt}}(s, a) & \text{probability } 1 - \epsilon \\ \text{random from Actions}(s) & \text{probability } \epsilon \end{cases}$$

Function Approx. :Large State Spaces, hard to explore

Define features $\phi(s, a)$ and weights \mathbf{w} :

$$\hat{Q}_{\text{opt}}(s, a; \mathbf{w}) = \mathbf{w} \cdot \phi(s, a)$$

On each (s, a, r, s') :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta [\underbrace{\hat{Q}_{\text{opt}}(s, a; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}}] \phi(s, a)$$

Adversarial Games: search state plus these:

Utility(s): agent's utility for end state s

Player(s) \in Players: player who controls state s

Game Evaluation Recurrence(Generic)

$$V_{\text{eval}}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \sum_{a \in \text{Actions}(s)} \pi_{\text{agent}}(s, a) V_{\text{eval}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \sum_{a \in \text{Actions}(s)} \pi_{\text{opp}}(s, a) V_{\text{eval}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \end{cases}$$

Expectimax Recurrence(opponent takes average)

$$V_{\text{exptmax}}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} V_{\text{exptmax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \sum_{a \in \text{Actions}(s)} \pi_{\text{opp}}(s, a) V_{\text{exptmax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \end{cases}$$

Minimax Recurrence

$$V_{\minimax}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} V_{\minimax}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \min_{a \in \text{Actions}(s)} V_{\minimax}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \end{cases}$$

Relationship between game values

$$V(\text{Ex}, \pi_{i7}) \geq V(\max, \pi_{i7}) \geq V(\max, \pi_{\min}) \geq V(\text{Ex}, \pi_{\min})$$

Expectiminimax Algorithm

$$V_{\text{expectiminimax}}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} V_{\text{expectiminimax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \min_{a \in \text{Actions}(s)} V_{\text{expectiminimax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \\ \sum_{c \in \text{Actions}(s)} \pi_{\text{coin}}(s, c) V_{\text{expectiminimax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{coin} \end{cases}$$


Speeding up minimax: Eval func and Alpha Beta pruning

Depth Limited Search

$$V_{\minimax}(s, d) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \text{Eval}(s) & d = 0 \\ \max_{a \in \text{Actions}(s)} V_{\minimax}(\text{Succ}(s, a), d) & \text{Player}(s) = \text{agent} \\ \min_{a \in \text{Actions}(s)} V_{\minimax}(\text{Succ}(s, a), d-1) & \text{Player}(s) = \text{opp} \end{cases}$$

AlphaBeta Pruning

alpha: lwr bound on val of max node (store $\alpha_s = \max_{s' \leq s} a_s$)
beta: uppr bound on val of min node (and $\beta_s = \min_{s' \leq s} b_s$)



TD Learning(On policy and need Succ states)

On each (s, a, r, s') :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta [\underbrace{\hat{V}_\pi(s; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_\pi(s'; \mathbf{w}))}_{\text{target}}] \nabla_{\mathbf{w}} \hat{V}_\pi(s; \mathbf{w})$$

$$V(s; \mathbf{w}) = \mathbf{w} \cdot \phi(s)$$

$$\nabla_{\mathbf{w}} V(s; \mathbf{w}) = \phi(s)$$

Simultaneous Games(Rock paper scissors)

For every simultaneous two-player zero-sum game with a finite number of actions:

$$\max_{\pi_A} \min_{\pi_B} V(\pi_A, \pi_B) = \min_{\pi_B} \max_{\pi_A} V(\pi_A, \pi_B),$$

where π_A, π_B range over mixed strategies.

Non Zero Sum Games(Prisoner's dilemma)

In any finite-player game with finite number of actions, there exists at least one Nash equilibrium.

A Nash equilibrium is (π_A^*, π_B^*) such that no player has an incentive to change his/her strategy:

$$V_A(\pi_A^*, \pi_B^*) \geq V_A(\pi_A, \pi_B^*) \text{ for all } \pi_A$$

$$V_B(\pi_A^*, \pi_B^*) \geq V_B(\pi_A^*, \pi_B) \text{ for all } \pi_B$$

Factor graph

Variables:

$$X = (X_1, \dots, X_n), \text{ where } X_i \in \text{Domain}_i$$

Factors:

$$f_1, \dots, f_m, \text{ with each } f_j(X) \geq 0$$

Scope of a factor f_j : set of variables it depends on.

Arity of f_j is the number of variables in the scope.

Unary factors (arity 1); Binary factors (arity 2).

Each assignment $x = (x_1, \dots, x_n)$ has a weight:

$$\text{Weight}(x) = \prod_{j=1}^m f_j(x)$$

$$\arg \max_x \text{Weight}(x)$$

A CSP is a factor graph where all factors are constraints:

$$f_j(x) \in \{0, 1\} \text{ for all } j = 1, \dots, m$$

The constraint is satisfied iff $f_j(x) = 1$.

An assignment x is consistent iff $\text{Weight}(x) = 1$ (i.e., all constraints are satisfied).

Backtrack($x, v, \text{Domains}$):

- If x is complete assignment: update best and return
- Choose unassigned VARIABLE X_i
- Order VALUES Domain_i of chosen X_i

For each value v in that order:

$$\delta \leftarrow \prod_{f_j \in D(x, X_i)} f_j(x \cup \{X_i : v\})$$

If $\delta = 0$: continue

Domains_i ← Domains via LOOKAHEAD

Backtrack($x \cup \{X_i : v\}, w\delta, \text{Domains}$)

Let $D(x, X_i)$ be set of factors depending on X_i and x but not on unassigned variables.

Forward Checking(Look Ahead)

After assigning a variable X_i , eliminate inconsistent values from the domains of X_i 's neighbors.

If any domain becomes empty, don't recurse.

When unassign X_i , restore neighbors' domains.

Least constrained value: useful when all factors are constraints (all assignment weights are 1 or 0)

Most constrained variable: useful when some factors are constraints (can prune assignments with weight 0)

Forward checking: need to actually prune domains to make heuristics useful!

Arc Consistency: eliminate values from domains-> reduce branching

A variable X_i is arc consistent with respect to X_j if for each $x_i \in \text{Domain}_i$, there exists $x_j \in \text{Domain}_j$ such that $f(\{X_i : x_i, X_j : x_j\}) \neq 0$ for all factors f whose scope contains X_i and X_j .
EnforceArcConsistency(X_i, X_j): Remove values from Domain_i to make X_i arc consistent with respect to X_j .

AC 3 Algorithm

Add X_j to set.

While set is non-empty:

- Remove any X_k from set.
- For all neighbors X_l of X_k :
 - Enforce arc consistency on X_l w.r.t. X_k .
 - If Domain_l changed, add X_l to set.

Runtime $O(E \Delta^3)$. Δ -> largest domain E -> number of edges


N ARY constraints

Variables: $X_1, X_2, X_3, X_4 \in \{0, 1\}$
Factor: $[X_1 \vee X_2 \vee X_3 \vee X_4]$

Auxiliary Variables hold intermediate computation.

Initialization: $[A_0] = 0$ i 0 1 2 3 4
Processing: $[A_i = A_{i-1} \vee X_i]$ X_i : 0 1 0 1
Final output: $[A_4] = 1$ A_i : 0 0 1 1 1

pack A_{i-1} and A_i into one variable B_i



Initialization: $[B_1[1] = 0]$

Processing: $[B_i[2] = B_i[1] \vee X_i]$

Final output: $[B_4[2] = 1]$

Consistency: $[B_{i-1}[2] = B_i[1]]$

Variables: B_i is (pre, post) pair from processing X_i

Pruning Techniques only useful for constraints, does not help reduce actual runtime.

Finding Maximum Weight Assignments: Greedy Algorithm

Partial assignment $x \leftarrow \{\}$

For each $i = 1, \dots, n$:

Extend: Compute weight of each $x_v = x \cup \{X_i : v\}$

Prune: $x \leftarrow x_v$ with highest weight

Beam Search Algorithm:

Initialize $C \leftarrow \{\}$


For each $i = 1, \dots, n$:

Extend: $C' \leftarrow \{x \cup \{X_i : v\} : x \in C, v \in \text{Domain}_i\}$

Prune: $C' \leftarrow K$ elements of C' with highest weights

Running time: $O(n(Kb) \log(Kb))$ with branching factor $b = |\text{Domain}|$, beam size K

Iterated conditional modes



$$o_1(x_1) f_1(x_1, v) o_2(v) f_2(v, x_3) o_3(x_3) \quad x_1, v, x_3$$

Initialize x to a random complete assignment

Loop through $i = 1, \dots, n$ until convergence:

Compute weight of $x_v = x \cup \{X_i : v\}$ for each v

$x \leftarrow x_v$ with highest weight

Gibbs Sampling: Sample an assign. Prop. to its weight

Note: The Gibbs update depends on the Markov blanket on a variable

Initialize x to a random complete assignment

Loop through $i = 1, \dots, n$ until convergence:

Compute weight of $x_v = x \cup \{X_i : v\}$ for each v

Choose $x \leftarrow x_v$ with probability prop. to its weight

Algorithms for max-weight assignments in factor graphs:

Backtracking search: exact, exponential time

Beam search: approximate, linear time

Iterated conditional modes: approximate, deterministic

Gibbs sampling: approximate, randomized

Conditioning Algorithm

To condition on a variable $X_i = v$, consider all factors f_1, \dots, f_k that depend on X_i .

Remove X_i and f_1, \dots, f_k .

Add $g_j(x) = f_j(x \cup \{X_i : v\})$ for $j = 1, \dots, k$.

Conditionally Independent: conditioning on C , makes A and B independent

Let $A \subseteq X$ be a subset of variables.

Define MarkovBlanket(A) be the neighbors of A that are not in A .

Elimination Algorithm

To eliminate a variable X_i , consider all factors f_1, \dots, f_k that depend on X_i .

Remove X_i and f_1, \dots, f_k .

$$\text{Add } f_{\text{new}}(x) = \max_{x_i} \prod_{j=1}^k f_j(x)$$

Running time: $O(n \cdot |\text{Domain}|^{\max\text{-arity}+1})$

The treewidth of a factor graph is the maximum arity of any factor created by variable elimination with the best variable ordering.

Bayesian Ntwks: Joint, Marginal(sum), Conditional(select rows) Distribution: three types for random variables

Explaining away: Suppose two causes positively influence an effect. Conditioned on the effect, conditioning on one cause reduces the probability of the other cause

Let $X = (X_1, \dots, X_n)$ be random variables.

A Bayesian network is a directed acyclic graph (DAG) that specifies a joint distribution over X as a product of local conditional distributions, one for each node:

$$\mathbb{P}(X_1 = x_1, \dots, X_n = x_n) \stackrel{\text{def}}{=} \prod_{i=1}^n p(x_i | x_{\text{Parents}(i)})$$

All factors (local conditional distributions) satisfy:

$$\sum_{x_i} p(x_i | x_{\text{Parents}(i)}) = 1 \text{ for each } x_{\text{Parents}(i)}$$

Probabilistic Inference Strategy: $\mathbb{P}(Q | E = e)$

- Remove (marginalize) variables that are not ancestors of Q or E .
- Convert Bayesian network to factor graph.
- Condition on $E = e$ (shade nodes + disconnect).
- Remove (marginalize) nodes disconnected from Q .
- Run probabilistic inference algorithm (manual, variable elimination, Gibbs sampling, particle filtering).

HMM:

$$\mathbb{P}(H = h, E = e) = \underbrace{p(h_1)}_{\text{start}} \prod_{i=2}^n \underbrace{p(h_i | h_{i-1})}_{\text{transition}} \prod_{i=1}^n \underbrace{p(e_i | h_i)}_{\text{emission}}$$

Smoothing Query: Lattice Representation:

Forward: $F_i(h_i) = \sum_{h_{i-1}} F_{i-1}(h_{i-1}) w(h_{i-1}, h_i)$

Backward: $B_i(h_i) = \sum_{h_{i+1}} B_{i+1}(h_{i+1}) w(h_i, h_{i+1})$

Define $S_i(h_i) = F_i(h_i) B_i(h_i)$

Gibbs Sampling

Loop through $i = 1, \dots, n$ until convergence:

Set $X_i = v$ with prob. $\mathbb{P}(X_i = v | X_{-i} = x_{-i})$

(notation: $X_{-i} = X \setminus \{X_i\}$)

Particle Filtering

Initialize $C \leftarrow \{\}$

For each $i = 1, \dots, n$:

Propose (extend): $C' \leftarrow \{x \cup \{X_i : x_i\} : x \in C, x_i \sim p(x_i | x_{i-1})\}$

Reweight:

Compute weights $w(x) = p(e_i | x_i)$ for $x \in C'$

Resample (prune):

$C \leftarrow K$ elements drawn independently from $\omega \propto w(x)$

Max Likelihood Learning Algorithm

Input: training examples D_{train} of full assignments

Output: parameters $\theta = \{p_d : d \in D\}$

Count:

For each $x \in D_{\text{train}}$:

For each variable x_i :

Increment $\text{count}_{d_i}(x_{\text{Parents}(i)}, x_i)$

Normalize:

For each d and local assignment $x_{\text{Parents}(i)}$:

Set $p_d(x_i | x_{\text{Parents}(i)}) \propto \text{count}_{d_i}(x_{\text{Parents}(i)}, x_i)$

data = {(d, d), (d, d), (d, d), (c, 1), (c, 5)}

Regularization:Laplace Smoothing(add 6 for Dice)

just add to the count for each possible value, regardless of whether it was observed or not. Eg: $X = 6$ for dice

$$P_{\text{LAP},k}(x) = \frac{c(x) + k}{N + k|X|}$$

Expectation maximization (EM)algorithm


E-step:

- Compute $q(h) = \mathbb{P}(H = h | E = e; \theta)$ for each h (use any probabilistic inference algorithm)
- Create weighted points: (h, e) with weight $q(h)$

M-step:

- Compute maximum likelihood (just count and normalize) to get θ

Repeat until convergence.



Estimate Likely No of Heads and Tails for First Toss

For A - $H = .45 * 5 = 2.2$, $T = .45 * 5 = 2.2$

For B - $H = .55 * 5 = 2.8$, $T = .55 * 5 = 2.8$

Max Likelihood of $p^n(1-p)^m = n/(n+m)$

E step: for each missing assume all vals, create virtual bags

and find probs. Mstep: find Max likelihood for all params.